



# A /goal for *Claude Code*

Codex CLI ships with one. Claude Code didn't.

*So I built it. May 2026.*



# The kit

*Portable. Drop into any Claude Code project.*

**GOAL.md schema**

Objective · Non-goals · Validation · Budget · Stop conditions · Progress log

**SessionStart hook**

Auto-injects the active goal into every Claude Code session

**Five slash commands**

/goal · /goal-status · /goal-pause · /goal-resume · /goal-clear

**Python loop runner**

Per-run JSON archive · USD / iteration / token caps

**Bash variant**

Parity with the Python runner

**SDK harness stub**

Tighter runtime continuation via the Agent SDK

# Overnight test

*Built a Jeopardy! training kit from public prep literature. Just for me.*

**8 iterations · 7 documents · ~\$6 total · agent self-terminated**

Status: **done**. Agent self-flipped to terminal state — exactly the behavior we wanted from a long-running loop.

All 7 deliverables landed. Total ~150KB of content across the kit:

|                             |           |       |                          |
|-----------------------------|-----------|-------|--------------------------|
| prep-strategies-research.md | 208 lines | 18 KB | (iter 1, manual)         |
| training-system.md          | 291 lines | 22 KB | (iter 3, with directive) |
| resources.md                | 258 lines | 19 KB | (iter 6, overnight)      |
| progress-tracker.md         | 155 lines | 10 KB | (iter 7, overnight)      |
| drills/buzzer-drills.md     | 218 lines | 16 KB | (iter 2, manual)         |
| drills/knowledge-drills.md  | 221 lines | 20 KB | (iter 4, overnight)      |
| drills/wagering-drills.md   | 405 lines | 26 KB | (iter 5, overnight)      |

The consistency pass caught a real bug. This is the highlight. The agent found that `progress-tracker.md` defined "LBR" as "Last-Beat Reps" while `buzzer-drills.md` and `training-system.md` defined it as "Lost Buzzer Races" — different metrics with opposite higher-is-better polarity. Aligned tracker to match. That's exactly the kind of cross-document drift that a final consistency pass is supposed to catch, and it did.

Non-goals all held. `Project_Pipeline.md` mtime unchanged from yesterday morning. Nothing in `goal-kit/` was touched. Zero permission\_denials across all overnight runs.

Cost estimate. I can only see the last iteration's JSON (\$0.60 for the consistency pass), but extrapolating from the pattern: ~\$0.55-0.60/iter overnight × 4 iterations + \$0.60 consistency pass + \$3.14 already spent = ~\$6.00 total for the full goal. Well under what I'd budgeted.

Loop discipline. 8 progress log entries total, newest-first, all in the format we specified. Agent followed the cap (would have stopped at 12 anyway). Token usage stayed inside the 250K cumulative bound.

*The consistency pass caught three files defining the same acronym with opposite polarity.*

# What the build surfaced

*Three things I won't run a long-horizon loop without.*

1

## Non-goals over step prompts

The Non-goals section did more work than the step prompt itself. Telling the agent what NOT to touch was the load-bearing constraint.

2

## Per-deliverable validation

Validation rules per deliverable are the difference between trusting the loop and re-reading every output by hand.

3

## Final consistency pass

Non-optional. It catches what individual iterations cannot — like the same acronym defined three different ways.